

## Lecture 10

By - Amal Sunny

---

# Classical Games

## Motivation for Classical Games

- Simple rules, deep concepts
- Studied and played for hundred/thousands of years
- Good IQ test for AI
- Chess was the most basic case study for RL/AI
- And... well games are fun 😊

## State of the Art: AI in games

Program	Level of Play	Program to Achieve Level
Checkers	Perfect	<i>Chinook</i>
Chess	Superhuman	<i>Deep Blue</i>
Othello	Superhuman	<i>Logistello</i>
Backgammon	Superhuman	<i>TD-Gammon</i>
Scrabble	Superhuman	<i>Maven</i>
Go	Grandmaster	<i>MoGo<sup>1</sup>, Crazy Stone<sup>2</sup>, Zen<sup>3</sup></i>
Poker <sup>4</sup>	Superhuman	<i>Polaris</i>

---

<sup>1</sup>9 × 9

<sup>2</sup>9 × 9 and 19 × 19

<sup>3</sup>19 × 19

<sup>4</sup>Heads-up Limit Texas Hold'em

- Superhuman - beat the best human player
- Perfect - Solved the game
- Grandmaster/Master/International Master - levels of
- This is for all AI methods, next we showcase how far RL has gone for those games.

## State of the Art: RL in games

Program	Level of Play	RL Program to Achieve Level
Checkers	<b>Perfect</b>	<i>Chinook</i>
Chess	International Master	<i>KnightCap / Meep</i>
Othello	<b>Superhuman</b>	<i>Logistello</i>
Backgammon	<b>Superhuman</b>	<i>TD-Gammon</i>
Scrabble	<b>Superhuman</b>	<i>Maven</i>
Go	<b>Grandmaster</b>	<i>MoGo<sup>1</sup>, Crazy Stone<sup>2</sup>, Zen<sup>3</sup></i>
Poker <sup>4</sup>	Superhuman	<i>SmooCT</i>

---

<sup>1</sup>9 × 9

<sup>2</sup>9 × 9 and 19 × 19

<sup>3</sup>19 × 19

<sup>4</sup>Heads-up Limit Texas Hold'em

---

## Game Theory - Optimality in Games

---

- In a multiplayer game, we can't just have a simple fixed policy as each player's action depends very much on what the other player does.
- Our question is to find the optimal policy  $\pi^i$  for the  $i$  th player
- So to make this an RL problem, we fix the policy of all other players as  $\pi^{-i}$  and consider it part of the environment for the  $i$  th player.
- But we still need to find the best way to play the game regardless of whatever opponent we have, i.e. generally best policy.
- For that we first define Best response  $\pi_*^i(\pi^{-i})$  which is the optimal policy against those strategies  $\pi^{-i}$
- And then we use Nash equilibrium as the joint policy for all players.
  - Nash's equilibrium states that every player playing this game would pick the most optimal one bearing in mind other player's strategies
  - i.e. every player's policy is the best response
  - no one would deviate from it

$$\pi^i = \pi_*^i(\pi^{-i})$$

- This might not be the most optimal for *all* opponents, but it acts as the best one against any general opponent.

## Single-Agent and Self-play RL

- Best response is solution to single-agent RL problem
  - where we include other players as part of the environment, factoring in their policies into our

## MDP

- Game is reduced to an MDP
- Solve for the MDP to get optimal response
- Nash equilibrium is fixed-point of self-play RL
  - We generate experience by playing games between 2 agents we made following the same policy
  - Each agent learns the best response to other player, based on the experience tuples generated.
  - And each agent's environment changes based on each player's policy improving
  - i.e. both players are adapting to each other
  - And we do this iterative process until we reach a point, where no further improvements are being made - the Nash equilibrium (if it exists for the problem)

## Two-player Zero-Sum Games

- We focus on a special class of games:
  - Two player games - have two (alternating players)
    - P1 is taken as white and P2 as black
  - Zero sum game - Has equal and opposite rewards for Black and White
    - $R^1 + R^2 = 0$
- We will consider the following methods to find Nash equilibria in these games
  - Game tree search(i.e. planning)
  - Self-play RL
- We further divide the games into 2 classifications - Perfect, Imperfect information games
  - Perfect information (or Markov) games are fully observed
    - Chess
    - Checkers
    - Othello
    - Backgammon - is stochastic, but all states are observed and known
    - Go
  - Imperfect information ones are partially observed
    - Scrabble - dont know opponents hand
    - Poker - same as above
- We focus on perfect information ones first

## Minimax

---

- A value function defines the expected total reward given joint policies  $\pi = \langle \pi^1, \pi^2 \rangle$

$$v_{\pi}(s) = E_{\pi}[G_t | S_t = s]$$

- A minimax value function maximizes white's expected return while minimizing black's expected return
  - i.e. it assumes each player plays alternatively to estimate value function at every point assuming players played ideally ahead


$$v_*(s) = \max_{\pi^1} \min_{\pi^2} v_{\pi}(s)$$

- A minimax policy is a joint policy  $\pi = \langle \pi^1, \pi^2 \rangle$  that achieves the minimax values. We use a tree like search to evaluate that
  - For larger trees, alpha-beta search is used which cuts away parts of the tree that are not of interest.
- There is a unique minimax value function
- A minimax policy is a Nash equilibrium


## Minimax Search


- Minimax values are evaluated by depth-first game tree search

Example



- Initially we're given value functions of the bottom most nodes.
- We must find minimax policy with this





- However this tree would grow exponentially for actual games
- Impractical to search till the end - thus we use value function approximator  $v(s, w) \approx v_*(s)$ 
  - also called evaluation function, heuristic function
- Value function is used to estimate minimax value at leaf nodes
- And search is ran to fixed depth w.r.t to leaf values.
- We'll consider an example of this - Deep Blue

## Deep Blue

- First actual chess program to beat a top level human player
- Knowledge
  - 8000 handcrafted chess features
  - Binary-linear value function
  - Weights were largely tuned by human experts
- Search
  - High performance parallel alpha-beta search
  - 480 special-purpose VLSI chess processors
  - Searched 200 million positions/second
  - Looked ahead 16-40 ply
- Results
  - Defeated human champion Garry Kasparov 4-2 (1997)
  - Most watched event in internet history
- But this was not RL, this was completely handcrafted features combined with minimax searching.

## Chinook

- Is a Checkers program.
- Knowledge
  - Binary-linear value function
  - 21 knowledge-based features (position, mobility, ...)
  - x4 phases of the game
- Search
  - High performance alpha-beta search
  - Retrograde analysis
    - Search backward from won positions
    - Store all winning positions in lookup tables
    - Plays perfectly from last n checkers
- Results

- Defeated Marion Tinsley in world championship 1994
  - won 2 games but Tinsley withdrew for health reasons
- Chinook solved Checkers in 2007
  - perfect play against God

## Self-Play TD Learning

- We apply previously learnt RL methods to RL algorithms involving games with self-play
- MC: update value function towards the return  $G$
- TD(0): update value function towards successor value  $v(S_{t+1})$
- TD( $\lambda$ ): update value function towards the  $\lambda$ -return  $G_t^\lambda$

## Policy Improvement with Afterstates

- For deterministic games, we need only to estimate  $v_*(s)$  for policy selection.
- This is because we can efficiently evaluate the afterstate

$$q_*(s, a) = v_*(succ(s, a))$$

- Rules of the game would deterministically give us the successor state  $succ(s, a)$ .
- Actions are then selected - min/max-ing depending on black or white's turn.
- This improves joint policy for both.

## Self-Play TD in Othello: Logistello

- Created its own features
  - Started with raw input features eg: black stone at c1
  - Constructed new features by conjunction/disjunction
  - Created 1.5m features in diff configs
  - Binary-linear value function using these features
- Logistello used generalized policy iteration
  - Generate batch of self-play games from current policy
  - Evaluate policies using Monte-Carlo (regress to outcomes)
  - Greedy policy improvement to generate new players
- Results
  - Defeated World Champion Takeshi Murakami 6-0

## Self-Play TD in Backgammon: TD-Gammon

- The board is flattened into a row of stripes, and features are extracted out of the position of beads on those rows - using a neural network.
- Initialized with random weights
- Trained by games of self-play
- Using non-linear TD learning


- Greedy policy improvement (no exploration)
  - Exploration not required due to the stochastic nature of the game (die roll)
- Algorithm always converged in practice
- Not true for other games
  - That is due to the dice roll again - the randomness smooths out the value function.
- Results
  - Zero expert knowledge - strong intermediate play
  - Hand-crafted features  $\Rightarrow$  advanced level of play (1991)
  - 2-ply search  $\Rightarrow$  strong master play (1993)
  - 3-ply search  $\Rightarrow$  superhuman play (1998)
  - Defeated world champion Luigi Villa 7-1 (1992)

## Combining RL and Minimax Search

---

### Simple TD

- TD: updates value towards successor value



- Value function approximator  $v(s, w)$  with parameter  $w$
- Value function gets backed up from raw value at next state

$$v(S_t, w) \leftarrow v(S_{t+1}, w)$$

- First learn the entire value function by TD learning
- Then use value function in minimax search(no learning)


$$v_+(S_t, w) = \text{minimax}_{\text{leaves}(S_t)} v(s, w)$$

Results:

- Othello: superhuman performance in Logistello
- Backgammon: superhuman performance in TD-Gammon
- Chess: poor performance
- Checkers: poor performance
- In chess tactics seem necessary to find signal in position
  - i.e. we need to use search

## TD root

- TD root: update value towards successor search value



- Search value is computed at root position  $S_t$

$$v_+(S_t, w) = \text{minimax}_{\text{seleaves}(S_t)} v(s, w)$$


- Value function based up from search value at next state

$$v(S_t, w) \leftarrow v_+(S_{t+1}, w) = v(l_+(S_{t+1}), w)$$

- where  $l_+(s)$  is the leaf node achieving minimax value from s

## TD Leaf

- TD leaf: update search value towards successor search value



- Search value computed at current and next step

$$v_+(S_t, w) = \text{minimax}_{\text{seleaves}(S_t)} v(s, w), \quad v_+(S_{t+1}, w) = \text{minimax}_{\text{seleaves}(S_{t+1})} v(s, w)$$

- Search value at step t backed up from search value at t+1.

$$\begin{aligned} v_+(S_t, w) &\leftarrow v_+(S_{t+1}, w) \\ \implies v(l_+(S_t), w) &\leftarrow v(l_+(S_{t+1}), w) \end{aligned}$$

## TD leaf in Chess: Knightcap

- Learning
  - Knightcap trained against expert opponent
  - Starting from standard piece values only
  - Learnt weights using TD leaf




- Search
  - Alpha-beta search with standard enhancements
- Results
  - Achieved master level play after a small number of games
  - Was not effective in self-play
  - Was not effective without starting from good weights

## TD leaf in Checkers: Chinook

- Original Chinook used hand-tuned weights
- Later version was trained by self-play
- Using TD leaf to adjust weights
  - Except material weights which were kept fixed
- Self-play weights performed  $\geq$  hand-tuned weights
- i.e. learning to play at superhuman level

## TreeStrap

- TreeStrap: Updates search value towards deeper search values



- Minimax search value computed at *all* nodes  $s \in \text{nodes } (S_t)$
- Value backed up from search value, at same step, for all nodes

$$v(s, w) \leftarrow v_+(s, w)$$

$$\Rightarrow v(s, w) \leftarrow v(l_+(s), w)$$

- TreeStrap uses all the info from our search tree instead of just taking one - doesn't waste the information gained in the search, uses it all.

## TreeStrap in Chess: Meep

- Binary linear value function with 2000 features
- Starting from random initial weights (no prior knowledge)
- Weights adjusted by TreeStrap
- Won 13/15 vs. international masters
- Effective in self-play
- Effective from random initial weights

# Simulation-Based Search

- Self-play RL can replace search
- Instead of searching, we just simulate how a game would have gone with whatever policies we've formulated used for both players from current root state  $S_t$
- Apply RL to simulated experience
  - MC control => MC tree search
  - Most effective variant is UCT algorithm
    - Balances exploration/exploitation in each node using UCB
  - Self-play UCT converges on minimax values

## Performance of MCTS (MC Tree Search) in Games

- MCTS is best performing method in many challenging games
  - Go
  - Hex
  - Lines of Action
  - Amazons
- In many games simple Monte-Carlo search is enough
  - Scrabble
  - Backgammon

## Simple MC search in Maven

- Maven is a program built to play Scrabble
  - while normally we assume a computer would be better, due to having the entire dictionary in data - human players end up memorizing that + strategizing better
- Learning
  - Maven evaluates moves by score + v (rack)
  - Binary-linear value function of rack
  - Using one, two and three letter features
    - Q??????, QU?????, III????
  - Learnt by Monte-Carlo policy iteration (cf. Logistello)
- Search
  - Roll-out moves by imagining n steps of self-play
  - Evaluate resulting position by score + v (rack)
  - Score move by average evaluation in rollouts
  - Select and play highest scoring move
  - Specialised endgame search using B\*
- Results
  - Maven beat world champion Adam Logan 9-5
  - Analysis showed Maven had error rate of 3 points per game

## Smooth UCT search

- Apply MCTS to information-state game tree
- Variant of UCT, inspired by game-theoretic Fictitious Play
  - Agents learn against and respond to opponents average behaviour
- Extract average strategy from nodes action counts,

$$\pi_{avg}(a|s) = \frac{N(s, a)}{N(s)}$$

- At each node, pick actions according to

$$A \sim \begin{cases} UCT(S), & \text{with probability } \mu \\ \pi_{avg}(\cdot|S), & \text{with probability } 1 - \mu \end{cases}$$

- Empirically, in variants of Poker:
  - Naive MCTS diverged
  - Smooth UCT converged to Nash equilibrium

## RL in games: Successful Recipe

Program	Input features	Value Fn	RL	Training	Search
Chess <i>Meep</i>	Binary <i>Pieces, pawns, ...</i>	Linear	TreeStrap	Self-Play / Expert	$\alpha\beta$
Checkers <i>Chinook</i>	Binary <i>Pieces, ...</i>	Linear	TD leaf	Self-Play	$\alpha\beta$
Othello <i>Logistello</i>	Binary <i>Disc configs</i>	Linear	MC	Self-Play	$\alpha\beta$
Backgammon <i>TD Gammon</i>	Binary <i>Num checkers</i>	Neural network	TD( $\lambda$ )	Self-Play	$\alpha\beta$ / MC
Go <i>MoGo</i>	Binary <i>Stone patterns</i>	Linear	TD	Self-Play	MCTS
Scrabble <i>Maven</i>	Binary <i>Letters on rack</i>	Linear	MC	Self-Play	MC search
Limit Hold'em <i>SmooCT</i>	Binary <i>Card abstraction</i>	Linear	MCTS	Self-Play	-